

CS 143 Notes

Indexes

Sparse – (key, ptr) per block, not every tuple has a (key,pointer) pair in the index

Dense – every tuple has a (key, ptr) pair in the index

Sparse/Primary Index/Clustering

tuples ordered according to search key

(key, pointer) to first tuple in "block" but not to every tuple in

table

Secondary/Non-clustering Index

Tuples in table are not ordered by index search key

Must always have a first level dense index

Sparse index from second level and above

– clustered index: rows stored physically on disk in same order as index

– non clustered index: second list that has pointers to physical rows, can

have many non clustered indexes

btree insertion

On Leaf insert overflow:

Split overflowing node

copy first key of newly split node to parent non-leaf node

On NonLeaf insert overflow

Split overflowing node

Move the key up into parent

A BTreeNode should have at least

Leaf:

$\text{ceil}((n+1)/2)$ pointers

$\text{ceil}(n+1)/2 - 1$ keys;

NonLeaf:

$\text{ceiling}(n/2)$ pointers

$\text{ceiling}(n/2) - 1$ keys

Root

2 pointers, 1 key

btree deletion

BTree Deletion

Leaf

No Underflow

Delete value from leaf

Coalesce leaf with neighbor

node underflow, neighboring node not filled up yet

Move keys from node with underflow into neighboring node (R → L)

Delete that node from the tree

Delete the key of the deleted node in the parent node

Redistribute leaf with neighbor

node that is filled up, another underflowed when deleted from

=> overflow caused when merging neighboring nodes

Nonleaf

- Coalesce non-leaf with neighbor

 - after deleting from leaf and from parent, parent underflows

 - Delete the parent node that is underflowing

 - Pull down the midkey of the parent's parent, merge midkey with rest of merged sibling non-leaf node

- Redistribute non-leaf with neighbor

 - after deleting from leaf and from parent, parent underflows

 - Pull down mid key of the parent of underflowing node to merge

 - with sibling non-leaf AND IT OVERFLOWS

 - do overflow algo by moving new sibling key to parent

If a node is deleted delete the pointer and the adjacent key (to the right) from the above parent node.

Extendible Hash Insertion:

- If a hashbucket overflows

 - If (overflow hashbucket i == directory i)

 - double directory size copy existing pointers

 - increase directory i by 1

 - split overflowing bucket

 - update directory pointers

 - increase child hashbuckets by 1

Extendible Hash Merge (after deletion)

- Bucket i 's must be the same

- First $(i-1)$ bits must be the same

Directory Shrink

- All bucket i 's are smaller than directory i

ER diagram

ER Relationship Notes

Entity -> Rectangle

Relationship -> Diamond

Total Participation -> double Line (entity participates in relationship at least once)

General Cardinality Notation

Along the edge of a relationship

min..max

the entity participates min to max times

- arrow points at the "one" end

Multiple entities can participate in one relationship

eg(Students and TA take classes)

Roles

If a entity participates multiple times in a relationship you can label each edge as a different role

(eg students are partners and partners can take different roles)

Subclasses

Student isa foreign or domestic student

Use a triangle along with subclassing

Subclass participates in the relationships of its superclass

Subclass may participate in its own relationship

Subclass inherits all traits of superclass

Generalization: subclass \rightarrow superclass

Specialization: superclass \rightarrow subclass

Total specialization: Entity is always one of subclasses (student always either domestic or foreign) \rightarrow double line

Weak Entity Sets (WES)

entity sets without keys

Double Rectangle and Double Diamond

Discriminator: set of attributes in WES that are part of the key
– dashed underling

Owner Entity Set (OES): entity set providing part of the key

Identifying Relationship: Link between WES and OES

Converting E/R to Tables

Strong Entity Set: A table with all attributes

Relationship Set: Table contains keys of linked entities and own attributes of relationship

WES: Table with its own attributes and keys from owner ES

Do not need to translate identifying relationship

Subclass

1. Table for each subclass with attributes and superclass key
2. A big Table with all attributes with null for missing ones

Normalization theory

Redundancies in table result in the following anomalies:

Update anomaly

Insertion anomaly

Deletion anomaly

Functional Dependency

– each X value in R is associated with precisely one Y value in R

$X \rightarrow Y$

u_1, u_2 in R if $u_1[X] = u_2[X]$ then $u_1[Y] = u_2[Y]$

Trivial FD $X \rightarrow Y$ and Y contained in X (sid,name \rightarrow sid)

Nontrivial FD: $X \rightarrow Y$ Y not contained in X (sid \rightarrow name, sid, addr)

Completely non-trivial FD: $X \rightarrow Y$ but Y has no overlap between X and Y (sid \rightarrow name, addr)

Key

X is a key of R iff X^+ results in all attributes in R and no smaller subset fulfills the first requirement

Lossless Decomposition occurs if
 $R(X,Y,Z) \rightarrow R_1(X,Y) \text{ or } R(X,Z)$
 if $X \rightarrow Y$ or $X \rightarrow Z$

Shared attributes is the key of one of the tables

////////////////////////////////////

BCNF Form

 R is in BCNF iff every non-trivial $X \rightarrow Y$, X contains a key
 no redundancy from FD

BCNF Normalization Algorithm

 For any R in schema

 If non-trivial $X \rightarrow Y$ holds in R and X does not have the key

 Compute X^+

 Split R in $R(X^+)$ and $R(X,Y)$ where Y is elements in R except X^+

ex:

ClassInstructor(dept, cnum, title, unit, instructor, office, fax)

instructor \rightarrow office

office \rightarrow fax

dept, cnum \rightarrow title, unit

dept, cnum \rightarrow instructor

$R_1(\text{inst, office, fax}) \Rightarrow R_3(\text{office, fax}) R_4(\text{inst, office})$

$R_2(\text{dept, cnum, title, unit, inst})$

////////////////////////////////////

Multi-Valued Dependency and 4NF

$X \twoheadrightarrow Y$

 if for tuples u and v, $u[X] = v[X]$ then there exists a tuple w s.t

$w[X] = u[X] = v[X]$

$w[Y] = u[Y]$

$w[Z] = v[Z]$

– for all X's that are the same with different Y's, the Z's can be swapped

ex: course \twoheadrightarrow book, course \twoheadrightarrow lecturer

Course	Book	Lecturer
AHA	Silberschatz	John D
AHA	Nederpelt	John D
AHA	Silberschatz	William M
AHA	Nederpelt	William M
AHA	Silberschatz	Christian G
AHA	Nederpelt	Christian G
OS0	Silberschatz	John D
OS0	Silberschatz	William M

Complementation Rule

$X \twoheadrightarrow Y$ then $X \twoheadrightarrow Z$ where Z is all attributes in R except X and Y

Trivial MVD

$X \twoheadrightarrow Y$ is trivial if

 Y in X

$X \cup Y = R$

Fourth Normal Form (4NF)

R is in 4NF if for every nontrivial FD $X \rightarrow Y$ or MVD $X \twoheadrightarrow Y$, X contains a key
no redundancy from FD or MVD
if 4NF, is BCNF

4NF Decomposition

For any R in the schema

If nontrivial $X \twoheadrightarrow Y$ holds on R and if X does not contain a key

Split R in $R_1(X, Y)$ and $R_2(X, Z)$ where Z is all attributes in R except

(X, Y)

YUS68

ex:

Class(dept, cnum, title, ta, sid, name)

dept, cnum \rightarrow title

sid \rightarrow sname

dept, cnum \rightarrow ta

$R_1(\text{sid}, \text{sname})$

$R_2(\text{dept}, \text{cnum}, \text{title}, \text{ta}, \text{sid})$

$\Rightarrow R_3(\text{dept}, \text{cnum}, \text{title})$

$R_4(\text{dept}, \text{cnum}, \text{ta}, \text{sid})$

$\Rightarrow R_5(\text{dept}, \text{cnum}, \text{ta})$

$R_6(\text{dept}, \text{cnum}, \text{sid})$

If $X \rightarrow Y$ then $X \twoheadrightarrow Y$

Simplified 4NF: R is 4NF if for every nontrivial MVD $X \twoheadrightarrow Y$, X contains a key

Transaction

ACID

Atomicity

Consistency

Isolation

Durability

Serial Schedule: all operations in any transaction are performed without any interleaving

Conflict Serializable Schedule:

Non-Conflicting action: A pair of actions that do not change the result when we swap them

Conflicting (two) actions are categorized by

Involve the same objects/variables

At least one of the two actions is a write

Conflict Equivalence

S1 is conflict Equivalent to S2 if S1 can be rearranged into S2 by a series of swaps of non-conflicting actions

Conflict Serializability

S1 is conflict serializable if it is conflict equivalent to some serial schedule

Precedence Graph P(S)

P(S) is acyclic \iff S is conflict serializable

You can only check whether a schedule is conflict serializable not if it is conflict-equivalence

Algo:

- for each action in S
 - get all actions that affect same current action
 - draw arrow if at least 1 of pair is a write

Recoverable Schedule

S is Recoverable if T_j reads a data item written by T_i , the commit operation of T_i appears before commit of T_j

Dirty Read: reading a value that has been committed

Cascading rollback: A single transaction abort leads to a series of transaction rollback

Cascadeless Schedule:

S is cascadeless if T_j reads a data item written by T_i , the commit operations appears before the READ of T_j

Cascadeless \rightarrow Recoverable

Recoverable DOES NOT imply cascadeless

Relationship Between Schedules

Recoverable contains Cascadeless contains Serializable but Conflict Serializable (around serial) not contained in any of those

We want either Conflict Serializable + Recoverable or Conflict serializable + Cascadeless

////////////////////////////////////

Locking Protocol

Rigorous Two Phase Locking Protocol (R2PL)

Rule(1): T_i has to lock tuple tk before any read/write

Rule(2): When T_i is holding the lock on tk , T_j cannot obtain the lock on tk for ($j \neq i$)

Rule(3): Release all locks at commit

Rigorous 2PL ensures a conflict-serializable and cascadeless schedule
commit vs time graph: staircase up, at commit, # locks drops to 0

abruptly

Two Phase Locking Protocol (2PL)

Rule(1): T_i lock a tuple before any read/write

Rule(2): If T_i holds the lock on A. T_j cannot access A ($j \neq i$)

Rule(3): Two stages

(a): growing stage: T_i may obtain locks but may not release any lock

(b): shrinking stage: T_i may release lcks, but may not obtain any lock

2PL ensures a conflict-serializable schedule

commit vs time graph: staircase up then down

Conflict Serializable contains 2P,L, 2PL contains R2PL

Recovery and Logging

Log Based Recovery

Log record written to disk before the actual data is written to disk

Before commit of T_i all log records for T_i written to disk including commit

During abort DBMS gets old values from the log

During recovery, DBMS does the following

executes all actions in the log from the beginning

rolls back all actions of "non-committed" transactions in reverse order

Use checkpoint to minimize recovery time

SQL Isolation Levels

Isolation Level	dirty read	nonrepeatable read	phantom
read uncommitted	Y	Y	Y
read committed	N	Y	Y
repeatable read	N	N	Y
serializable	N	N	N

Nonrepeatable read: when T_i reads the same row multiple times, T_i may get different values

Phantom: When new tuples are inserted, some of them are seen by statements

* SET AUTOCOMMIT {0|1}

* SET TRANSACTION READ ONLY, ISOLATION LEVEL REPEATABLE READ

Only when all transactions are serializable, we guarantee ACID

Read only transactions help DBMS optimize for more concurrency

Join Algos

Nested Loop Join

```
For each r in R
  for each s in S
    if r.C = s.C then output r,s
COST:  $b_R + \text{ceil}(b_R / (M - 2)) * b_S$ 
- good for small relations
```

Index Join

```
for each r in R
  x <- index look up in S (S.c, r.c)
  for each s in X
    write matched tuples
```

-> |R|
-> C
-> J

COST: $b_R + |R| * (C + J)$
C - avg index lookup cost
J - # matching tuples
- good if index exists already

Sort merge

```
Sort stage: sort R and S
COST (to sort 1 table R):  $2b_R * (\text{ceil}(\log_{M-1}(b_R / M)) + 1)$ 
- don't need output buffer
Join stage: 2 ptr technique to join based on equality
COST:  $b_R + b_S$ 
- need output buffer
- good for non-equi-join
- sorted run: blocks / memory buffer
```

Hash join

```
Hashing stage:
  hash R tuples into  $G_k$  buckets
  hash S tuples into  $H_k$  buckets
join stage:
  for i = 1 to k
    match tuples in  $G_i, H_i$  buckets
COST:  $3(b_R + b_S)$ 
- best for equi-join and if relations not sorted and have no index
```